

COMP 408

Advanced Topics in Artificial Intelligence

Lecture 1

Introduction and Regular Expressions

8 / 2 / 2025

Some Slides are by D. Jurafsky and J. M. Martin

Course Information

- **Course Code:** COMP 408
- **Course Title:** Advanced Topics in Artificial Intelligence:
 - Part 1: **N**atural **L**anguage **P**rocessing (**NLP**)
 - Part 2: **M**ulti **A**gent **S**ystem (**MAS**) Dr. Hewayda
- **Number of Credit Hours:** 3 (3 hours Lecture + 0 Lab.)
- **Prerequisite :** -
- **Final Exam (duration):** 3 hours
- **Total Marks:** 150
 - 105 Final Exam.
 - 37 Mid-Term Exam.
 - 8 Oral Exam.

Part 1: NLP


Objectives

- Learn the broad topics, fundamental concepts, and some common techniques in the field of natural language processing.
- Understand the capabilities, limitations, and promise of NLP.
- Understanding of the computational properties of natural languages and the **commonly used algorithms** for processing linguistic information.

Textbook

Daniel Jurafsky and James H. Martin (2025), **Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition with Language Model**, 3rd edition, Pearson Education Limited.

<https://web.stanford.edu/~jurafsky/slp3/>



Natural Languages

- One of the fundamental aspects of human behavior and is a crucial component of our lives.
- A **communication mechanism** whose medium is **text or speech**.

Natural Language Processing (NLP)

- A field of computer science and artificial intelligence that focused on the **technology and algorithms of processing natural languages**.

Why processing natural languages?

- There are two main reasons to process natural languages:
 - **Enable machines to communicate with humans.**
 - Alan Turing proposed his Test based it on language.
 - **Acquire information from written language.**
 - Text summarization.
 - Sentiment analysis.
 - Search for relevant text document.
 - Classify text documents.
 - Topic modeling; identify the topic of a given text.

Common Applications of NLP

- Spell Checking
- Sentiment Analysis
- Information Retrieval
- Machine Translation
- Question Answering
- Email Classification

Categories of Knowledge

Types of knowledge required to deal with natural language:

- Phonology
- Morphology
- Syntax
- Semantics
- Pragmatics

Phonetics

- Study of the **speech sounds** that make up languages.

Morphology

- **Word** formation and their meaning.
- How words can be constructed from more basic meaning units called **morphemes**:
- A morpheme is the primitive unit of meaning in a language:
 - The meaning of the word "friendly" is derivable from the meaning of the noun "**friend**" and the suffix "**-ly**", which transforms a noun into an adjective.

English morphology

- English **affixes** can be divided to:
prefixes, suffixes
- For example:
 - **Eats** consists of the **stem** eat and the suffix -s
 - **Foxes** consists of the **stem** fox and the suffix -es
 - **Unstable** consists of the **stem** stable and the prefix un-

Arabic morphology

- Arabic **affixes** can be divided to:
prefixes, suffixes, infix
- For example:
 - **وا** suffix consists of the stem **كتب**
 - **ي** prefix consists of the stem **كتب**

Syntax

- How words can be put together to form **grammatical sentences**:
 - I saw a woman in the garden (**grammatical**)
 - * I **a** woman saw in the garden (**ungrammatical**)
- The statement “**I saw a man with a telescope**” is grammatical; has two parse trees **each of the trees corresponds to different meaning**.
 - This is called **syntactic ambiguity**.


Syntax Cont.

- A sentence is grammatical if it has a parse tree or derivation.
- Consider the context-free grammar:

$S \rightarrow NP \ VP$	$NP \rightarrow I$
$NP \rightarrow NP \ PP$	$P \rightarrow \text{with}$
$NP \rightarrow \text{Det } N$	$\text{Det} \rightarrow \text{the}$
$VP \rightarrow VP \ PP$	$N \rightarrow \text{man}$
$VP \rightarrow V \ NP$	$N \rightarrow \text{telescope}$
$PP \rightarrow P \ NP$	$V \rightarrow \text{saw}$

Syntax Cont.

Two parse trees can be constructed, for the sentence, from the pervious grammar



Semantics

- Study of the word **meaning** and how the meanings combine in sentences to form sentence meanings.
- This is the study of **context independent** meaning - the meaning a sentence has **regardless of the context** in which it is used:
 - The word **bank** has two different **senses** (meaning)
 - Ambiguity in the **sense** of the word

Pragmatics

- Study of the meaning of languages in **contexts**.

Question

"The tires are brand new."

- Given that the person **uttering** this sentence is responding to a complaint that the car is too cold, this sentence is:
 - A. Syntactically incorrect.
 - B. Semantically incorrect.
 - C. Pragmatically correct.
 - D. Syntactically and semantically correct.

Question

"The tires are brand new."

- Given that the person **uttering** this sentence is responding to a complaint that the car is too cold, this sentence is:
 - A. Syntactically incorrect .
 - B. Semantically incorrect.
 - C. Pragmatically correct.
 - D. Syntactically and semantically correct.**

What makes NLP hard?

- **Ambiguity:** Human language is full of ambiguities.
 - Words can have multiple meanings depending on context, which makes it hard for machines to understand the intended meaning.
- **Context:** Grasping the context is crucial.
 - Words and sentences are often interpreted based on the surrounding text, previous sentences, or even real-world knowledge.
- **Diversity of Languages:** Each language has its own set of grammar rules, syntax, idioms, and slang.
 - NLP models need to handle this diversity to understand and generate text accurately.

What makes NLP hard? Cont.

- **Idiomatic Expressions:** Phrases like "فات القطار" or "piece of cake" can be challenging since their meanings are not literal.
- **Sarcasm and Irony:** Detecting sarcasm or irony is tough because it relies heavily on tone, context, and cultural nuances, which are difficult for machines to grasp.
- **Named Entity Recognition (NER):** Identifying proper nouns, like names of people, places, or organizations, requires understanding of the text and sometimes external knowledge.
- **Morphology and Syntax:** Human languages have complex morphology (word forms) and syntax (sentence structures) that need to be parsed accurately.

Ambiguity Example

- Find at least 5 meanings of this sentence:
 - I made her duck

Ambiguity Example

- Find at least 5 meanings of this sentence:
 - I made her duck
1. I cooked waterfowl for her benefit (to eat)
 2. I cooked waterfowl belonging to her
 3. I created the (plaster?) duck she owns
 4. I caused her to quickly lower her head or body
 5. I waved my magic wand and turned her into undifferentiated waterfowl

Ambiguity Example

- I caused her to quickly lower her head or body
 - **Lexical category**: “duck” can be a **Noun** or a Verb
- I cooked waterfowl belonging to her.
 - **Lexical category**: “her” can be a possessive (“of her”) or dative (“for her”) pronoun
- I made the (plaster) duck statue she owns
 - **Lexical Semantics**: “make” can mean “create” or “cook”

Ambiguity

- **Grammar:** Make can be:
 - **Transitive: (verb has a noun direct object)**
 - I cooked [waterfowl belonging to her]
 - **Ditransitive: (verb has 2 noun objects)**
 - I made [her] (into) [undifferentiated waterfowl]
 - **Action-transitive (verb has a direct object and another verb)**
 - I caused [her] [to move her body]

Ambiguity

- **Phonetics!**
 - I mate or duck
 - I'm eight or duck
 - Eye maid; her duck
 - Aye mate, her duck
 - I maid her duck
 - I'm aid her duck
 - I mate her duck
 - I'm ate her duck
 - I'm ate or duck
 - I mate or duck

Example 1

I sent her messages

Has two different meanings

a. I sent some messages to her

b. I sent the messages written by her

- There is ambiguity in the word *her*, since it can be a **dative pronoun** or a **possessive pronoun**.
- There is also a syntactic ambiguity because the verb send is syntactically ambiguous, it may be transitive with one object as in (b) or with two objects as in (a).

Example 2

He drew one card

Has two different meanings

- a. He created one card having a picture.
- b. He chose one card from a set of cards.
- The different meanings are because of **semantic ambiguity** because the verb **draw** has many senses (meaning) such as **created a picture** and **choose**.

Regular Expression

Chapter 2 of the textbook

A tool for describing text patterns

Regular Expressions and Text Searching

- Regular expressions are a **compact textual representation** of a set of strings representing a language
 - In the simplest case, regular expressions describe **regular languages**
 - Here, a **formal language** means a set of strings from a given some alphabet.

Regular expressions

- A **formal language** for specifying text strings
- How can **we search for mentions** of these cute animals in text?
 - woodchuck
 - woodchucks
 - Woodchuck
 - Woodchucks
 - Groundhog
 - groundhogs



Regular Expressions: Disjunctions

- Letters inside square brackets [] specify **disjunction of characters**
- [wW] matches w or W

Pattern	Matches	Example
/[wW]oodchuck/	Woodchuck, woodchuck	<u>Woodchuck</u>
/[abc]/	'a', 'b', or 'c'	I like Java <u>a</u>
/[1234567890]/	Any digit from 0 to 9	<u>7</u> or 5

Regular Expressions: Disjunctions

- Ranges [A-Z]

Pattern	Matches	Example
/[A-Z]/	An upper case letter	<u>D</u> renched Blossoms
/[a-z]/	A lower-case letter	<u>m</u> y beans were impatient
/[0-9]/	A single digit	Chapter <u>1</u> : Down the Rabbit Hole
/[2-4]/	2, 3, or 4	Chapter <u>2</u>

Question

- The regular expression `B[0-37]` matches:
 1. `B37`
 2. `B3`
 3. `B4`
 4. `3`

Question

- The regular expression `B[0-37]` matches:

(`B[0, 1, 2, 3, 7]`)

1. `B37`

2. `B3`

3. `B4`

4. `3`

Regular Expression

`/B[0-37]/g` **2 matches**

Test String

`B37`
`B3`
`B4`
`3`

Regular Expression: More Disjunction

- The pipeline symbol ‘|’
- Examples,
 - $/a|b|c/ = /[abc]/$
 - $/[gG]roundhog/ \mid /[Ww]oodchuck/$

Regular expressions: ? * + .

? Marks **optionality** of the previous expression.

* Marks **zero or more** occurrences of the previous expression.

+ Marks **one or more** occurrences of the previous expression.

Period (.) Matches **any single** character.

Pattern	Matches	Example
/Colou?r/	previous character optional	<u>Color</u> <u>Colour</u>
/oo*h!/	Zero or more previous character	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> ...
/[ab]*/	Zero or more a's or b's	<u>aaa</u> , <u>abab</u> , <u>bbb</u>
to*	0 or more of previous char	<u>t</u> <u>to</u> <u>too</u> <u>tooo</u>
to+	1 or more of previous char	<u>to</u> <u>too</u> <u>tooo</u> <u>toooo</u>
/beg.n/		beg'n begin <u>begun</u>



Stephen C Kleene

Kleene *, Kleene +

Note

- The period (.) is often used with Kleene star (*), to mean any string of characters:
 - The regular expression `/cat.*cat/` matches any line with word cat at the beginning of the line and then any sequence of characters then the word cat.

Example

- How can we search for any of these?
 - woodchuck
 - woodchucks
 - Woodchuck
 - Woodchucks
 - Groundhog
 - groundhogs



Example

- How can we search for any of these?
 - woodchuck
 - woodchucks
 - Woodchuck
 - Woodchucks
 - Groundhogs
 - groundhogs
- `/[wW]oodchucks?| [Gg]roundhogs?/`**

A word begins with either w or W and may end with s or a word begins with G or g and may end with s



Regular expressions : Negation in Disjunction

Carat (^) as first character in [] negates the list

Note: Carat means negation only when it's first in []

Pattern	Matches	Examples
<code>/[^A-Z]/</code>	Not an upper-case letter	O <u>y</u> fn pripetchik
<code>/[^Ss]/</code>	Neither 'S' nor 's'	<u>I</u> have no exquisite reason"
<code>/[e^]/</code>	either 'e' or '^'	Look h <u>e</u> re
<code>/a^b/</code>	The pattern a caret b	Look up <u>a^b</u> now

Regular expressions: anchors `^`, `$`, `\b`, and `\B`

- The `^` matches the start of the line (if it appears at the beginning of the regular expression.)
- The `$` matches the end of the line.
- `\b` matches word boundary and `\B` matches non-word boundary

Pattern	Example
<code>/^The/</code>	<u>The</u> dog
<code>/^[^A-Za-z]/</code>	<u>1</u> “Hello”, the first <code>^</code> matches the start of line, the <code>^</code> inside <code>[]</code> means not.
<code>/cat\$/</code>	The big <u>cat</u> .

Note


- The caret, `^`, has three uses:
 1. To **match the start of the line**
 - `/^The/` matches the word **The** at the start of the line
 2. To indicate **negation** inside of square brackets
 - `/[^A-Z]/` matches non-uppercase letter
 3. To mean a caret
 - `/[ab^d]/` matches a, b, ^, or d
 - `/a^b/` matches `a^b`

Regular expressions: { }

Pattern	Match	Example
/the{3}/	3 occurrences of the last letter e	theee
/the{2, 4}/	From 2 to 4 occurrences of the last letter e	thee thee theeee
/the{2,}/	At least 2 occurrence of the last letter e	thee, thee, ...

RE operator precedence

- Parenthesis
- Counters
- Sequences and anchors
- Disjunction



Notes

1. Since counters have a higher precedence than sequence of characters:
 - the RE `/the*/` matches `theeee` not `thethe`.
2. Since sequence have a higher precedence than disjunction:
 - The RE `/the|any/` matches `the` or `any` not `theny`.

More operators

RE	Match
*	an asterisk “*”
\.	a period “.”
\?	A question mark
\d	[0-9], any digit
\D	[^0-9], any non-digit
\w	[a-zA-Z0-9_], any alphanumeric or underscore
\W	[^\w], a non-alphanumeric e.g., ? ; + * -
\s	[\r\t\n\f], white space
\S	[^\s], non-whitespace

Example 1

- Write a regular expression to match the word “the” (or “The”) in a text.
 - `/the/` misses capitalized T (The)
 - `/[tT]he/` incorrectly matches `others` or `theory`
 - `/\W[tT]he\W/`

False positives and false negatives

- The process we just went through was based on **two fixing kinds of errors**
 1. Not matching things that we should have matched (The)
False negatives
 2. Matching strings that we should not have matched (**there**, **then**, **other**)
False positives

Characterizing work on NLP

In NLP we are always dealing with these kinds of errors.

Reducing the error rate for an application often involves two antagonistic efforts:

- Increasing coverage (or *recall*) (minimizing false negatives).
- Increasing accuracy (or *precision*) (minimizing false positives)

All

only

Example 2

- Find regular expression to express the language with the following strings (**sheep language**):

baa! baaa! baaaa!

...

/baa*!/ or

/baa+!/

Example 3

- Write a regular expression to match the words `cat` and `dog` in a text:

– `/\b cat | dog \b /`

– `/[catdog]/` (not correct, why?)

- Write a regular expression to match both `guppy` and `guppies`:

– `/gupp (y|ies)/`

– `/guppy| ies` (not correct, why?)

<code>s r</code>	matches one string <code>s</code> or <code>r</code>
<code>[ab]</code>	matches one character <code>a</code> or <code>b</code>

Example 4

- A regular expression that matches 0 or more a's:

/a/*

- A regular expression that matches at least one a's:

/aa/*

Example 5

- A regular expression that matches a single digit
 - `/[0-9]/` or
 - `\d` (any digit)
- A regular expression that matches any non-digit
 - `\[^0-9]/` or
 - `\D`
- A regular expression that matches an integer:
 - `/[0-9][0-9]*/` or
 - `/[0-9]+/`

Example 6

- Give regular expression to represent the sets:

1. {lass, class, glass}

`/\b[cg]?lass\b/` or `/\b(c|g)? Lass\b/`

2. {jet, pet, net}

`/\b[jpn]et\b/` or

`/\b jet | pet | net \b/`

regexpal Software

← → ↻

Secure | https://www.regexpal.com


Dan's Tools
www.regexpal.com

Regex Pal

☰

[Ask a Question](#) [Related: Original Regexpal](#) [Donate!](#)

Regular Expression

JavaScript ▾  flags

`/[cg]?lass/g`

3 matches

Test String

`lass, class, glass`

Example 7

- Give regular expression to represent the sets:

1. {sun, sunday, sunrise, sunset}

`/sun(ε | day | rise | set)/`

or

`/sun(day|rise|set)?/`

Question

Select the correct regular expression that describes the set of all lower-case alphabetic strings with letter b as a second letter:

1. `/[a-z]*b[a-z]/`
2. `/[a-z]b[a-z]*/`
3. `/[a-z]+b[a-z]*/`
4. `/[a-z]*b[a-z]*/`